

Homework Set 1: COL380

Kartikeya Gupta, 2013CS10231

February 3, 2016

1.
 - In an un-pipelined processor, 1000 operations are processed and for each line, a total time of $2ns + 1ns + 1ns + 1ns + 1ns + 1ns + 2ns$ is required.
 \implies Total time needed = $1000 * 9ns$
 $\implies t = 9000ns$
 - In the given timings, the first and last stages take $2ns$ the amount of time while the other states take $1ns$ time. Because of this the time spent per instruction is $2ns$.
The number of stages in the pipeline are 7.
As there are 2 pipelined processors, the number of lines executed by each will be half of earlier = 500.
 \implies Time taken for processing 500 lines by a processor = $500 * 2ns = 1000ns$.
But we have not taken into account the time which the last instruction will spend before leaving the pipeline. This is going to be $1ns + 1ns + 1ns + 1ns + 1ns + 2ns = 7ns$.
 \implies Total time taken = $1000ns + 7ns = 1007ns$.

Diagram showing timeline:
TODO: Insert table
2.
 - For the peak operation, the entire y will be present in the cache and blocks of x will keep entering and getting evicted by the cache. The present z which is being needed will also be a member of the cache.
Let's now consider the memory access times as follows
 - (a) Time needed to get the entire y in cache from DRAM:
 K amount of cache lines to be retrieved = $K * 100ns$ time.
 - (b) Time needed to get the entire z in cache once from DRAM:
 K amount of cache lines to be retrieved = $K * 100ns$ time.
 - (c) Time needed to get the entire x in cache repeatedly from DRAM:
 $16K^2/4$ number of cache lines to be retrieved = $4K^2 * 100ns$
 - (d) Time needed for retrieving data from the caches for the arithmetic operations:

$16K^2$ operations taking place and for each of these 3 elements have to be accessed from the cache = $48K^2\text{ns}$

(e) Time needed for processing:

$2 * 16K^2$ operations require a time of = $32K^2\text{ns}$.

\Rightarrow Total time for this = $480K^2 + 200K\text{ns}$

Total number of instructions taking place = $32 * K^2$.

\Rightarrow Number of operations per second =

$$\frac{32K^2 * 10^9}{480K^2 + 200K}$$

Approximating K to 1000, \Rightarrow

$$\frac{32 * 10^6 * 10^3}{480 * 10^6 + 200 * 10^3} \text{Mflops}$$

\Rightarrow

$$66.68 \text{Mflops}$$

- Consider the following code for multiplying the matrices

```

1 for (int i=0; i< SIZE ; i++)
2     for (int j=0 ; j< SIZE ; j++)
3         for (int k=0; k< SIZE; k++)
4             C[i][j]+=A[i][k]*B[k][j];

```

For multiplying 2 dense matrices in given, we have to perform $(4K)^3$ mathematical operations

When the matrices are stored in row major form, for matrix A , we need to keep a given row i of the matrix A in the cache along with the memory in C where the result is to be stored. The elements of matrix B are to be fetched column wise but if we wish to access a particular element, we will get 3 other elements which are of no use at that time. Hence for a given i and j , the entire column from matrix C needs to be accessed from the DRAM.

Let us calculate the memory access times as follows:

(a) Time spent in getting Matrix A in the cache from DRAM:

$4K^2$ number of cache lines to be retrieved = $4K^2 * 100\text{ns}$ time.

(b) Time spent in getting Matrix C in the cache from DRAM:

$4K^2$ number of cache lines to be retrieved = $4K^2 * 100\text{ns}$ time.

(c) Time spent in getting Matrix B in the cache from DRAM:

$(4K)^3$ number of cache lines to be retrieved = $64K^3 * 100\text{ns}$ time.

(d) Time needed for retrieving data from caches for arithmetic operations:

$3 * (4K)^3$ values need to be accessed from the cache = $192K^3 \text{ ns}$ time.

(e) Time needed for the arithmetic operations to take place:
 $2 * (4K)^3$ number of arithmetic operations take place = $128K^3$ ns time.
 \Rightarrow Total time for this = $6720K^3 + 800K^2$ ns.
Total number of instructions taking place = $128K^3$
 \Rightarrow Number of operations per second =

$$\frac{128 * K^3 * 10^9}{6720 * K^3 + 800 * K^2} Flops$$

Approximating K to 1000 \Rightarrow

$$\frac{128 * 10^9 * 10^3}{6720 * 10^9 + 800 * 10^6} MFlops$$

\Rightarrow

$$19.045 MFlops$$

3.
 - The value assigned to y is 1. This is because when the instruction is executed on T_0 , the value of y is set to 1 and cache 2, snoops the change.
 - In the case when directory based cache coherence, the value is again 1. As x is in the shared section, when it gets updated by T_0 the changes are reflected.
 - There is no problem based on cache protocols and coherence in this situation. The problem here is because of a race condition.
4.
 -
 -
 -